

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
```

```
In [7]: plt.rcParams['figure.figsize'] = [19, 8]
warnings.filterwarnings('ignore')
```

```
In [8]: try:
cars_df = pd.read_csv("C:\\Users\\rakshith\\Downloads\\cars_df.csv",encoding="utf-8",nrows=20)
print("Dataset loaded successfully.")
except FileNotFoundError:
print("Error: 'cars.csv' not found. Please ensure the file is in the correct directory and update the path.")
if not cars_df.empty:
print("\n--- Initial Data Details ---")
print("Dataset shape (rows, columns):", {cars_df.shape})
print("\nColumn Info:")
cars_df.info()
```

Dataset loaded successfully.

--- Initial Data Details ---

Dataset shape (rows, columns): {(20, 15)}

Column Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 20 entries, 0 to 19

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Make	20 non-null	object
1	Model	20 non-null	object
2	Type	20 non-null	object
3	Origin	20 non-null	object
4	DriveTrain	20 non-null	object
5	MSRP	20 non-null	float64
6	Invoice	20 non-null	float64
7	EngineSize	20 non-null	float64
8	Cylinders	20 non-null	float64
9	Horsepower	20 non-null	float64
10	MPG_City	20 non-null	float64
11	MPG_Highway	20 non-null	float64
12	Weight	20 non-null	float64
13	Wheelbase	20 non-null	float64
14	Length	20 non-null	float64

dtypes: float64(10), object(5)

memory usage: 2.5+ KB

```
In [6]: print("\nFirst 5 rows of the dataset:")
print(cars_df.head())
```

First 5 rows of the dataset:

	Make	Model	Type	Origin	DriveTrain	MSRP	Invoice	\
0	Acura	MDX	SUV	Asia	All	36945.0	33337.0	
1	Acura	RSX Type S	2dr Sedan	Asia	Front	23820.0	21761.0	
2	Acura	TSX 4dr	Sedan	Asia	Front	26990.0	24647.0	
3	Acura	TL 4dr	Sedan	Asia	Front	33195.0	30299.0	
4	Acura	3.5 RL 4dr	Sedan	Asia	Front	43755.0	39014.0	

	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	\
0	3.5	6.0	265.0	17.0	23.0	4451.0	
1	2.0	4.0	200.0	24.0	31.0	2778.0	
2	2.4	4.0	200.0	22.0	29.0	3230.0	
3	3.2	6.0	270.0	20.0	28.0	3575.0	
4	3.5	6.0	225.0	18.0	24.0	3880.0	

	Wheelbase	Length
0	106.0	189.0
1	101.0	172.0
2	105.0	183.0
3	108.0	186.0
4	115.0	197.0

```
In [9]: print("\nLast 5 rows of the dataset:")
print(cars_df.tail())
```

Last 5 rows of the dataset:

	Make	Model	Type	Origin	DriveTrain	MSRP	\
15	Audi	A4 3.0 Quattro convertible 2dr	Sedan	Europe	All	44240.0	
16	Audi	A6 2.7 Turbo Quattro 4dr	Sedan	Europe	All	42840.0	
17	Audi	A6 4.2 Quattro 4dr	Sedan	Europe	All	49690.0	
18	Audi	A8 L Quattro 4dr	Sedan	Europe	All	69190.0	
19	Audi	S4 Quattro 4dr	Sedan	Europe	All	48040.0	

	Invoice	EngineSize	Cylinders	Horsepower	MPG_City	MPG_Highway	Weight	\
15	40075.0	3.0	6.0	220.0	18.0	25.0	4013.0	
16	38840.0	2.7	6.0	250.0	18.0	25.0	3836.0	
17	44936.0	4.2	8.0	300.0	17.0	24.0	4024.0	
18	64740.0	4.2	8.0	330.0	17.0	24.0	4399.0	
19	43556.0	4.2	8.0	340.0	14.0	20.0	3825.0	

	Wheelbase	Length
15	105.0	180.0
16	109.0	192.0
17	109.0	193.0
18	121.0	204.0
19	104.0	179.0

```
In [10]: print("\n--- Descriptive Statistics for Numerical Columns ---")
print(cars_df.describe())
```

```
--- Descriptive Statistics for Numerical Columns ---
```

	MSRP	Invoice	EngineSize	Cylinders	Horsepower	\
count	20.000000	20.000000	20.000000	20.000000	20.000000	
mean	41748.500000	37817.150000	3.060000	5.900000	238.750000	
std	15151.963838	13665.271774	0.703675	1.209611	47.040716	
min	23820.000000	21761.000000	1.800000	4.000000	170.000000	
25%	33371.250000	30349.250000	2.925000	6.000000	220.000000	
50%	38292.500000	34664.500000	3.000000	6.000000	220.000000	
75%	44705.000000	40331.250000	3.500000	6.000000	266.250000	
max	89765.000000	79978.000000	4.200000	8.000000	340.000000	

	MPG_City	MPG_Highway	Weight	Wheelbase	Length
count	20.000000	20.000000	20.000000	20.000000	20.000000
mean	18.900000	26.000000	3693.700000	107.100000	185.300000
std	2.44734	2.846974	403.228302	5.066921	8.657823
min	14.000000	20.000000	2778.000000	100.000000	172.000000
25%	17.000000	24.000000	3536.250000	104.000000	179.000000
50%	18.000000	25.000000	3726.000000	105.000000	181.500000
75%	20.000000	28.000000	3883.250000	109.000000	192.000000
max	24.000000	31.000000	4451.000000	121.000000	204.000000

```
In [13]: print("\nNull values before handling:")
cars_df = pd.read_csv("C:\\Users\\rakshith\\Downloads\\cars_df.csv", na_values=["NA", "Na", "?", "null", " "])
print(cars_df.isnull().sum())
cars_df.rename(columns={'MSRP': 'MRP', 'MPG_City': 'Mileage_City',
'MPG_Highway': 'Mileage_Highway'}, inplace=True)
print("\nColumns renamed. New column names:")
print(cars_df.columns)
print("\n--- Data Cleaning ---")
initial_rows = len(cars_df)
print(f"Number of duplicate rows found: {cars_df.duplicated().sum()}")
cars_df.drop_duplicates(inplace=True)
print(f"Duplicate rows dropped. Dataset shape is now: {cars_df.shape}")
```

Null values before handling:

```
Make      0
Model     0
Type      0
Origin    0
DriveTrain 0
MSRP      0
Invoice   0
EngineSize 0
Cylinders 2
Horsepower 0
MPG_City  0
MPG_Highway 0
Weight    0
Wheelbase 0
Length    0
dtype: int64
```

Columns renamed. New column names:

```
Index(['Make', 'Model', 'Type', 'Origin', 'DriveTrain', 'MRP', 'Invoice',
      'EngineSize', 'Cylinders', 'Horsepower', 'Mileage_City',
      'Mileage_Highway', 'Weight', 'Wheelbase', 'Length'],
      dtype='object')
```

--- Data Cleaning ---

Number of duplicate rows found: 0

Duplicate rows dropped. Dataset shape is now: (428, 15)

```
In [14]: if 'Cylinders' in cars_df.columns and cars_df['Cylinders'].isnull().any():
        median_cylinders = cars_df['Cylinders'].median()
        cars_df['Cylinders'].fillna(median_cylinders, inplace=True)
        print(f"\nMissing 'Cylinders' values filled with median value: {median_cylinders}")
        print("\nNull values after handling:")
        print(cars_df.isnull().sum())
        if 'Cylinders' in cars_df.columns:
            cars_df['Cylinders'] = cars_df['Cylinders'].astype(np.int64)
            print("\n'Cylinders' column data type converted to int64.")
            print(cars_df.dtypes)
```

Missing 'Cylinders' values filled with median value: 6.0

Null values after handling:

```
Make      0
Model     0
Type      0
Origin    0
DriveTrain 0
MRP       0
Invoice   0
EngineSize 0
Cylinders 0
Horsepower 0
Mileage_City 0
Mileage_Highway 0
Weight    0
Wheelbase 0
Length    0
dtype: int64
```

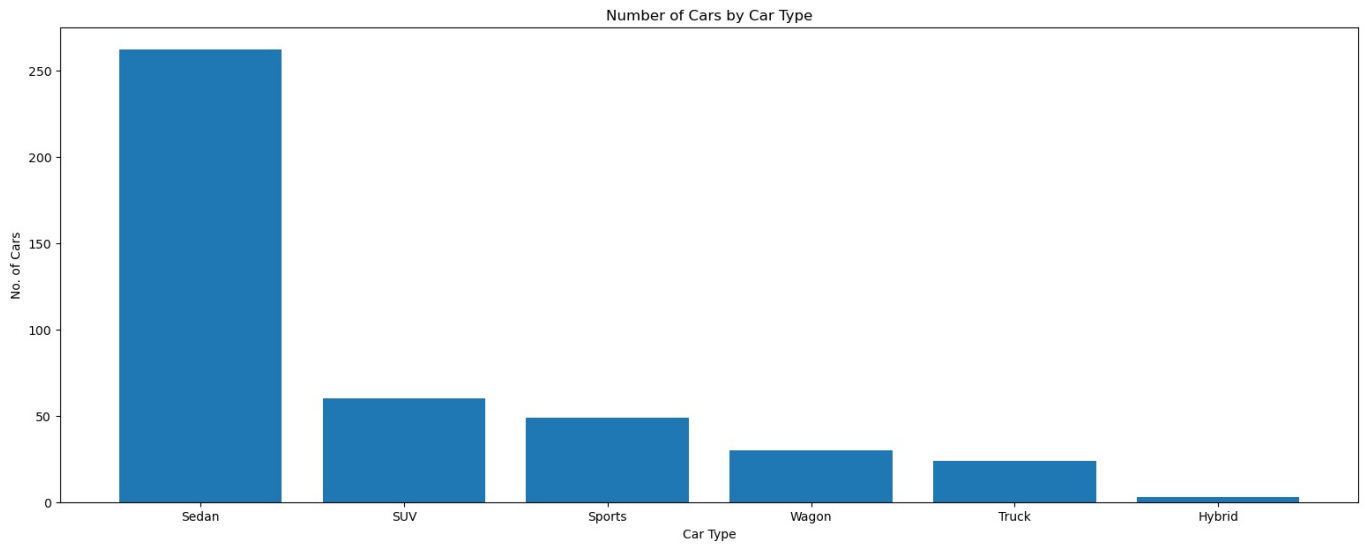
'Cylinders' column data type converted to int64.

```
Make      object
Model     object
Type      object
Origin    object
DriveTrain object
MRP       float64
Invoice   float64
EngineSize float64
Cylinders int64
Horsepower float64
Mileage_City float64
Mileage_Highway float64
Weight    float64
Wheelbase float64
Length    float64
dtype: object
```

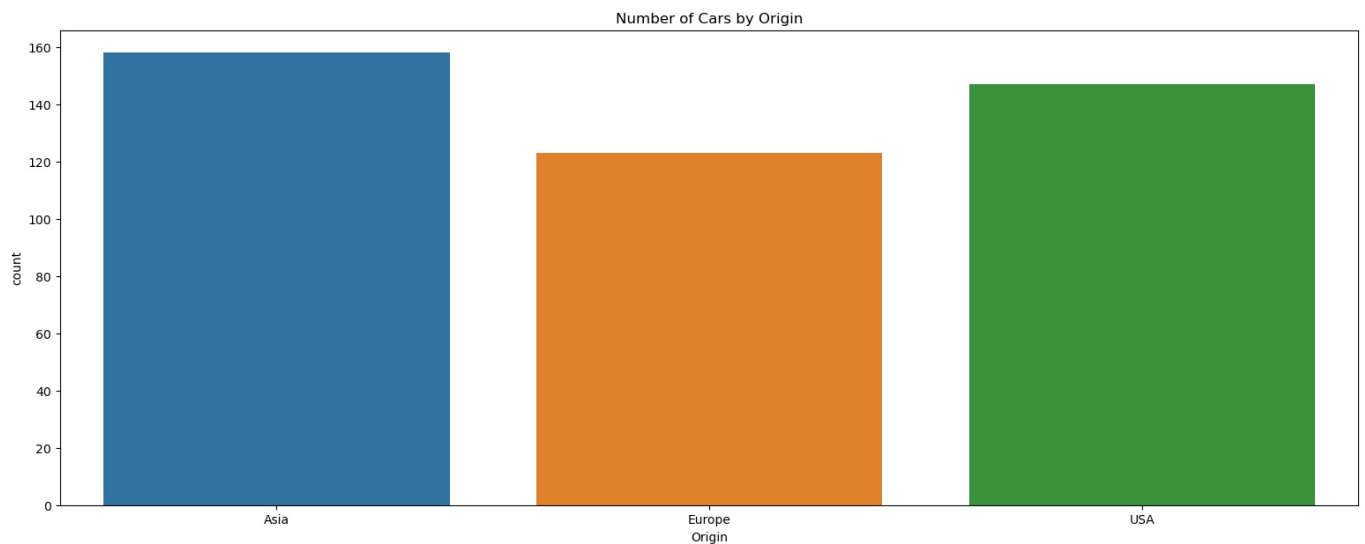
```
In [15]: print("\n--- Generating Visualizations ---")
        plt.figure(figsize=(19, 7))
        type_counts = cars_df['Type'].value_counts()
        plt.title("Number of Cars by Car Type")
        plt.bar(x=type_counts.index, height=type_counts.values)
        plt.xlabel("Car Type")
        plt.ylabel("No. of Cars")
```

```
plt.show()
```

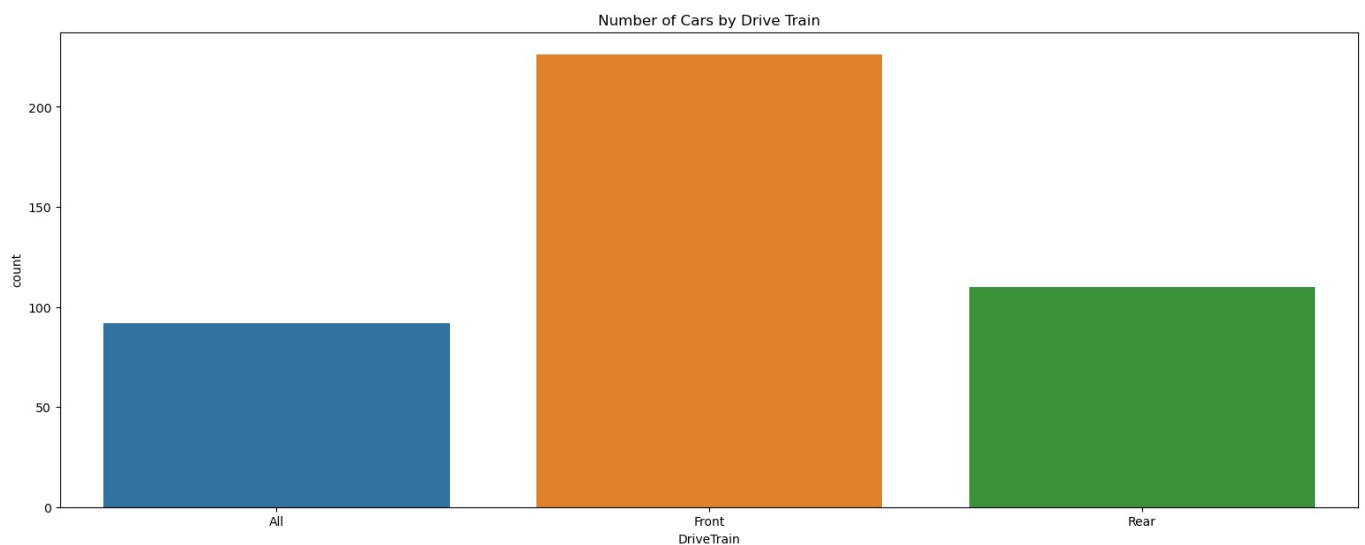
--- Generating Visualizations ---



```
In [16]: plt.figure(figsize=(19, 7))
plt.title("Number of Cars by Origin")
sns.countplot(data=cars_df, x="Origin")
plt.show()
```

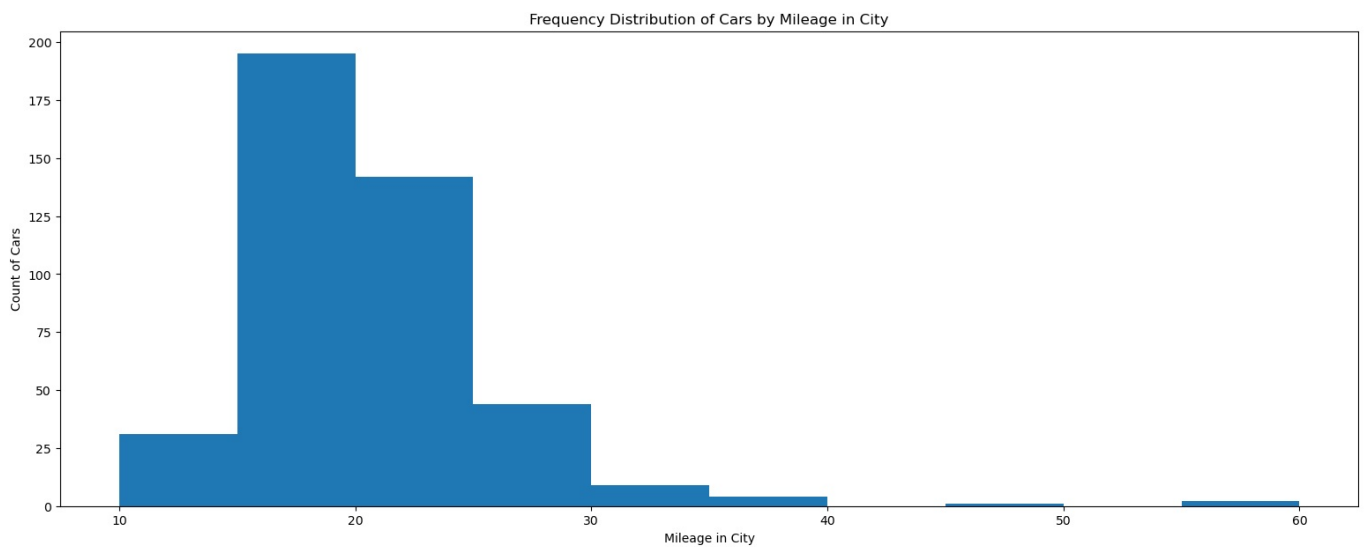


```
In [17]: plt.figure(figsize=(19, 7))
plt.title("Number of Cars by Drive Train")
sns.countplot(data=cars_df, x="DriveTrain")
plt.show()
```




```
In [18]: plt.figure(figsize=(19, 7))
plt.title("Frequency Distribution of Cars by Mileage in City")
plt.hist(x=cars_df['Mileage_City'])
plt.xlabel("Mileage in City")
plt.ylabel("Count of Cars")
```

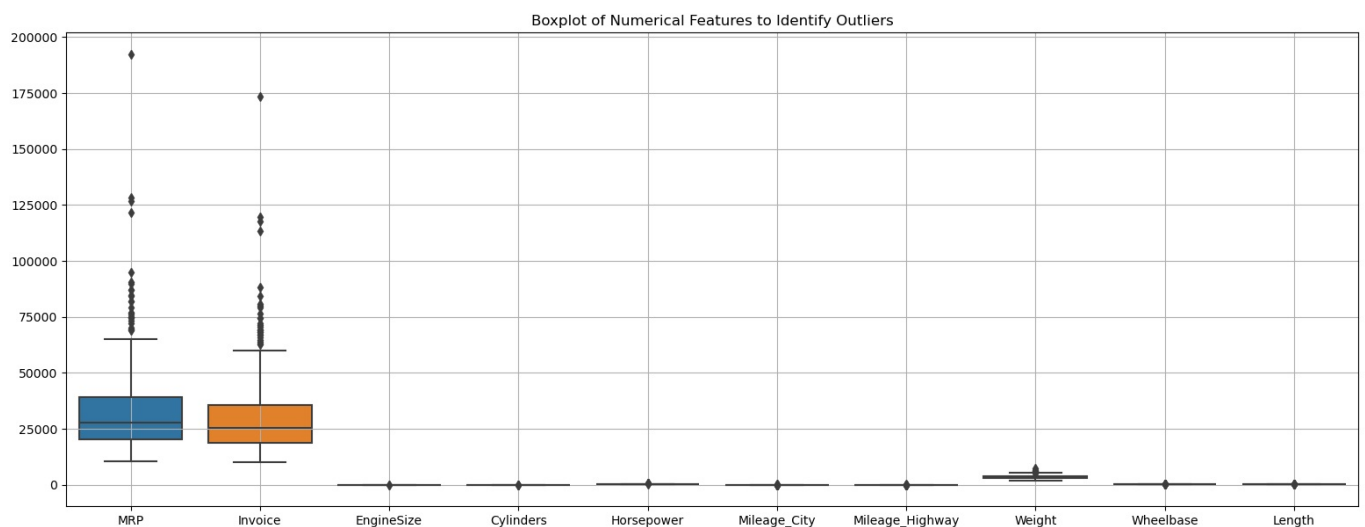
```
plt.show()
```



```
In [19]: plt.figure(figsize=(19, 7))
plt.title("Correlation Matrix of Numerical Features")
numerical_df = cars_df.select_dtypes(include=np.number)
sns.heatmap(numerical_df.corr(), annot=True)
plt.show()
```



```
In [20]: plt.figure(figsize=(19, 7))
plt.title("Boxplot of Numerical Features to Identify Outliers")
sns.boxplot(data=numerical_df)
plt.grid()
plt.show()
```



```
In [21]: print("\n--- Feature Engineering ---")
cars_cat_df = cars_df.select_dtypes(exclude=np.number)
cars_num_df = cars_df.select_dtypes(include=np.number)
print("Performing One-Hot Encoding on 'Type' column...")
```

```

from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder(sparse_output=False)
encoded_type = encoder.fit_transform(cars_cat_df[['Type']])
cars_encode_df = pd.DataFrame(encoded_type, columns=encoder.get_feature_names_out(['Type']))
print("One-Hot Encoded DataFrame head:")
print(cars_encode_df.head())

```

--- Feature Engineering ---

Performing One-Hot Encoding on 'Type' column...

One-Hot Encoded DataFrame head:

	Type_Hybrid	Type_SUV	Type_Sedan	Type_Sports	Type_Truck	Type_Wagon
0	0.0	1.0	0.0	0.0	0.0	0.0
1	0.0	0.0	1.0	0.0	0.0	0.0
2	0.0	0.0	1.0	0.0	0.0	0.0
3	0.0	0.0	1.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0	0.0

```

In [22]: print("\n--- Data Scaling ---")
print("Applying StandardScaler to numerical data...")
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
cars_num_scaled_df = pd.DataFrame(scaler.fit_transform(cars_num_df), columns=cars_num_df.columns)
print("Scaled numerical data description:")
print(cars_num_scaled_df.describe())

```

--- Data Scaling ---

Applying StandardScaler to numerical data...

Scaled numerical data description:

	MRP	Invoice	EngineSize	Cylinders	Horsepower
count	4.280000e+02	428.000000	4.280000e+02	4.280000e+02	4.280000e+02
mean	-1.577139e-16	0.000000	3.652322e-16	2.158191e-16	-7.470660e-17
std	1.001170e+00	1.001170	1.001170e+00	1.001170e+00	1.001170e+00
min	-1.158991e+00	-1.142905	-1.712933e+00	-1.808347e+00	-1.991379e+00
25%	-6.409709e-01	-0.632676	-7.421022e-01	-1.164443e+00	-7.091854e-01
50%	-2.648181e-01	-0.267866	-1.776656e-01	1.233648e-01	-8.202571e-02
75%	3.312970e-01	0.323216	6.351230e-01	1.233648e-01	5.451340e-01
max	8.227633e+00	8.146034	4.608757e+00	3.986788e+00	3.959670e+00

	Mileage_City	Mileage_Highway	Weight	Wheelbase	Length
count	4.280000e+02	4.280000e+02	4.280000e+02	4.280000e+02	4.280000e+02
mean	-2.448716e-16	-3.527811e-17	9.130806e-17	-6.474572e-16	2.573227e-16
std	1.001170e+00	1.001170e+00	1.001170e+00	1.001170e+00	1.001170e+00
min	-1.922891e+00	-2.588453e+00	-2.279333e+00	-2.307153e+00	-3.023605e+00
25%	-5.849947e-01	-4.958520e-01	-6.251890e-01	-6.208318e-01	-5.830854e-01
50%	-2.027386e-01	-1.470851e-01	-1.364646e-01	-1.390258e-01	4.447676e-02
75%	2.272995e-01	3.760652e-01	5.273695e-01	4.632317e-01	5.325807e-01
max	7.633511e+00	6.828252e+00	4.764630e+00	4.317679e+00	3.600663e+00

```

In [23]: cars_encode_df.reset_index(drop=True, inplace=True)
cars_num_scaled_df.reset_index(drop=True, inplace=True)
final_df = pd.concat([cars_encode_df, cars_num_scaled_df], axis=1)
print("\n--- Final Processed DataFrame ---")
print(f"Final shape: {final_df.shape}")
print(final_df.head())

```

--- Final Processed DataFrame ---

Final shape: (428, 16)

	Type_Hybrid	Type_SUV	Type_Sedan	Type_Sports	Type_Truck	Type_Wagon
0	0.0	1.0	0.0	0.0	0.0	0.0
1	0.0	0.0	1.0	0.0	0.0	0.0
2	0.0	0.0	1.0	0.0	0.0	0.0
3	0.0	0.0	1.0	0.0	0.0	0.0
4	0.0	0.0	1.0	0.0	0.0	0.0

	MRP	Invoice	EngineSize	Cylinders	Horsepower	Mileage_City
0	0.214856	0.188537	0.273884	0.123365	0.684503	-0.584995
1	-0.461376	-0.468388	-1.080764	-1.164443	-0.221395	0.752902
2	-0.298050	-0.304611	-0.719525	-1.164443	-0.221395	0.370645
3	0.021647	0.016134	0.002954	0.123365	0.754187	-0.011611
4	0.565724	0.510700	0.273884	0.123365	0.127028	-0.393867

	Mileage_Highway	Weight	Wheelbase	Length
0	-0.670235	1.151631	-0.259477	0.183935
1	0.724832	-1.055214	-0.861735	-1.001460
2	0.376065	-0.458983	-0.379929	-0.234440
3	0.201682	-0.003896	-0.018574	-0.025252
4	-0.495852	0.398428	0.824586	0.741768

```

In [24]: print("\n--- Outlier Treatment ---")
Q1 = final_df.quantile(0.25)
Q3 = final_df.quantile(0.75)
IQR = Q3 - Q1
lower_limit = Q1 - 1.5 * IQR
upper_limit = Q3 + 1.5 * IQR
rows_before = final_df.shape[0]

```

```

outliers_mask = ((final_df < lower_limit) | (final_df > upper_limit)).any(axis=1)
cars_df_no_outliers = final_df[~outliers_mask]
rows_after = cars_df_no_outliers.shape[0]
print(f"\nNumber of outliers detected and removed: {rows_before - rows_after}")
print(f"Shape of DataFrame after removing outliers: {cars_df_no_outliers.shape}")

```

--- Outlier Treatment ---


Number of outliers detected and removed: 203

Shape of DataFrame after removing outliers: (225, 16)

```

In [25]: plt.figure(figsize=(19, 7))
sns.boxplot(data=cars_df_no_outliers)
plt.title("Boxplot of Final Data After Outlier Removal")
plt.grid()
plt.show()
print("\nAnalysis complete.")

```



Analysis complete.

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js